

Homework 3

PSTAT 131

Contents

Binary Classification	1
---------------------------------	---

Binary Classification

For this assignment, we will be working with part of a Kaggle data set that was the subject of a machine learning competition and is often used for practicing ML models. The goal is classification; specifically, to predict which passengers would survive the Titanic shipwreck.

Load the data from `data/titanic.csv` into *R* and familiarize yourself with the variables it contains using the codebook (`data/titanic_codebook.txt`).

Notice that `survived` and `pclass` should be changed to factors. When changing `survived` to a factor, you may want to reorder the factor so that “Yes” is the first level.

Make sure you load the `tidyverse` and `tidymodels`!

Remember that you’ll need to set a seed at the beginning of the document to reproduce your results.

Question 1

Split the data, stratifying on the outcome variable, `survived`. You should choose the proportions to split the data into. Verify that the training and testing data sets have the appropriate number of observations. Take a look at the training data and note any potential issues, such as missing data.

```
# Load the data frist
titanic_data <- read_csv("titanic.csv")

# Convert 'survived' and 'pclass' to factors
titanic_data$survived <- as.factor(titanic_data$survived)
titanic_data$pclass <- as.factor(titanic_data$pclass)

# Set a seed for reproducibility
set.seed(42)

# Split the data, stratifying on 'survived'
data_split <- initial_split(titanic_data, prop = 0.7, strata = "survived")
train_data <- training(data_split)
test_data <- testing(data_split)

# Number of observations
cat("Number of observations in training set:", nrow(train_data), "\n")
```

```
## Number of observations in training set: 623
```

```
cat("Number of observations in testing set:", nrow(test_data), "\n")
```

```
## Number of observations in testing set: 268
```

```
# Check for missing data in the training set  
cat("Missing data in the training set:\n")
```

```
## Missing data in the training set:
```

```
print(summarize_all(train_data, funs(sum(is.na(.)))))
```

```
## # A tibble: 1 x 12  
##   passenger_id survived pclass  name  sex  age sib_sp parch ticket  fare cabin  
##         <int>    <int>  <int> <int> <int> <int> <int> <int>  <int> <int> <int>  
## 1             0        0      0     0    0  122      0     0      0     0  493  
## # i 1 more variable: embarked <int>
```

```
## From the table above, in the training set, there are 122 observations' "age"  
#is missing, and 493 observations' "cabin" is missing. And these missing data,  
#especially those without cabin info might negatively affect the accuracy of  
#the model.
```

Why is it a good idea to use stratified sampling for this data?

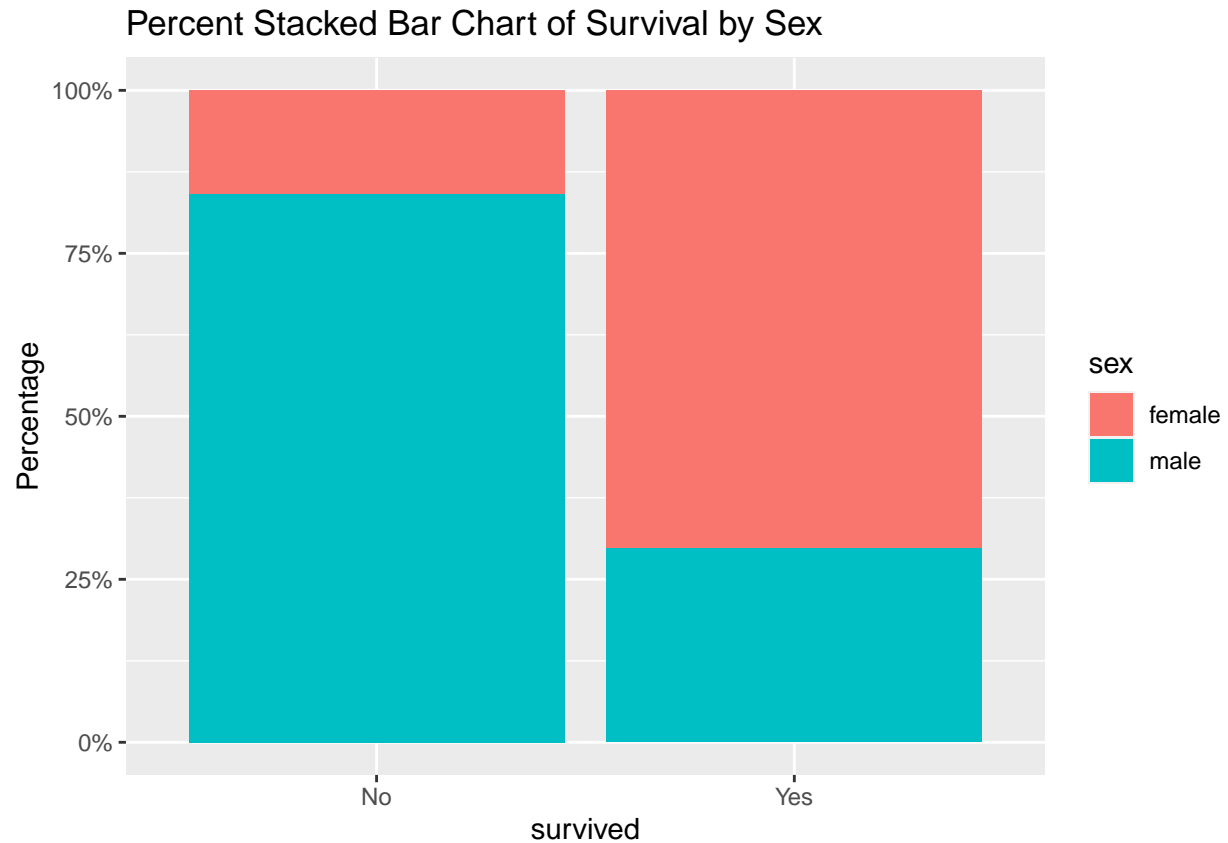
```
## In the context of the Titanic dataset, stratified sampling is particularly  
#important because the survival rate is historically low, leading to a class  
#imbalance in the data. By stratifying on the 'survived' column, the model is  
#exposed to a representative mix of survivors and non-survivors during both  
#training and testing. This approach mitigates the risk of the model being biased  
#towards predicting the majority class, thereby enhancing its predictive accuracy  
#specifically for this dataset.
```

Question 2

Using the **training** data set, explore/describe the distribution of the outcome variable **survived**.

Create a percent stacked bar chart (recommend using **ggplot**) with **survived** on the *x*-axis and **fill** = **sex**. Do you think **sex** will be a good predictor of the outcome?

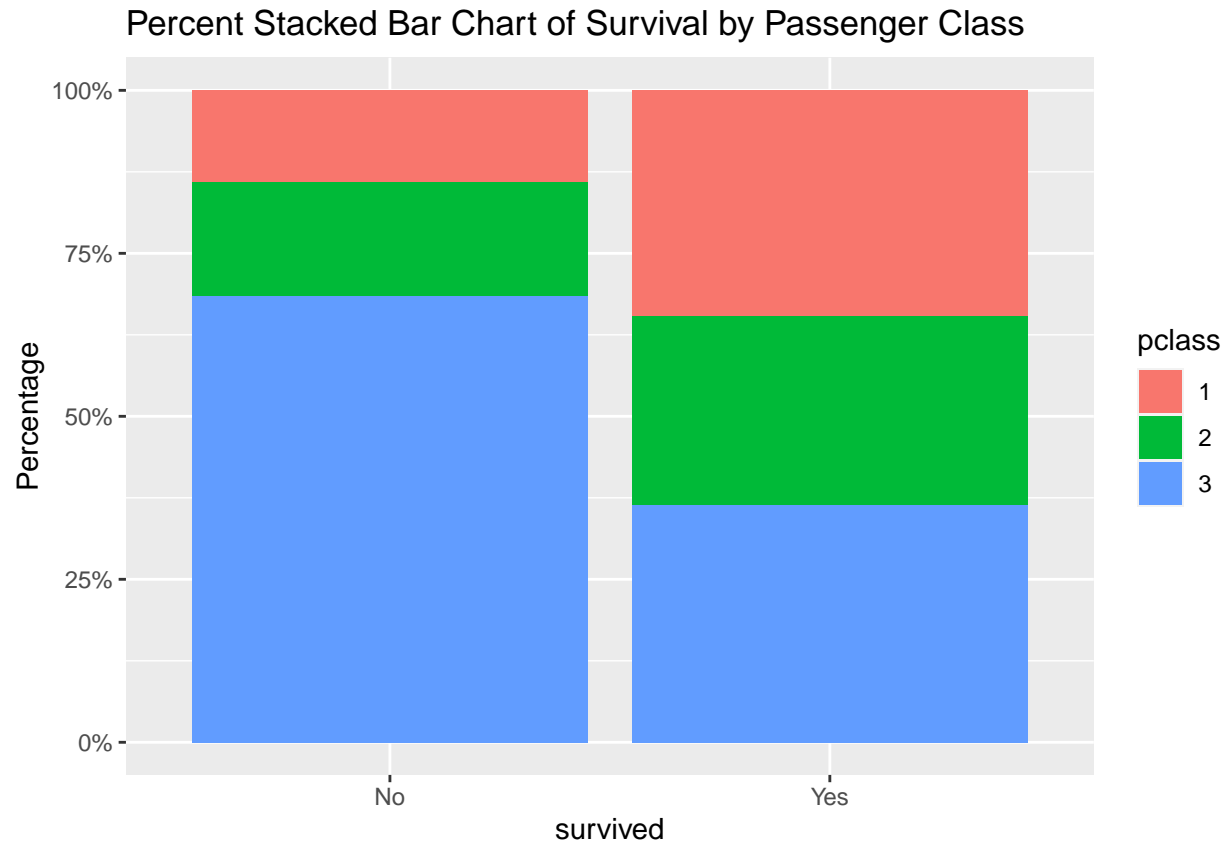
```
# Create the percent stacked bar chart per the requirement.  
ggplot(train_data, aes(x = survived, fill = sex)) +  
  geom_bar(position = "fill") +  
  scale_y_continuous(labels = scales::percent_format(scale = 100)) +  
  labs(title = "Percent Stacked Bar Chart of Survival by Sex",  
       y = "Percentage")
```



From the bar chart above, among those who did not survive, a significant majority are male (80%). And among those who did survive, a significant majority are female (70%). These substantial differences in survival rates between sex suggest that sex is likely to be a strong predictor of survival. The model would benefit from including this variable. So yes, it is a good predictor.

Create one more percent stacked bar chart of `survived`, this time with `fill = pclass`. Do you think passenger class will be a good predictor of the outcome?

```
# Create the percent stacked bar chart for 'pclass'
ggplot(train_data, aes(x = survived, fill = pclass)) +
  geom_bar(position = "fill") +
  scale_y_continuous(labels = scales::percent_format(scale = 100)) +
  labs(title = "Percent Stacked Bar Chart of Survival by Passenger Class",
       y = "Percentage")
```



From the bar chart above, Among those who did not survive, the majority were #from the third class with 65%, followed by 20% from the second class and 15% #from the first class. And among those who did survive, the distribution is more #even, with 36% from the first class, 36% from the third class, and 28% from the #second class. Specifically, passengers from the first class appear to have #higher chances of surviving compared to those from the third class, as evidenced #by the higher percentage of survivors in Pclass1 and the higher percentage of #non-survivors in Pclass3. So yes, 'pclass' will be a good predictor of the #outcome.

Why do you think it might be more useful to use a percent stacked bar chart as opposed to a traditional stacked bar chart?

The chart normalizes the data, making it easier to see that, for example, a #higher proportion of first-class passengers survived compared to third-class #passengers. This format is especially useful given that the number of passengers #varies between classes and genders, allowing for a more apples-to-apples #comparison. In other words, percent stacked bar chart makes it easier to read #the difference between multiple variables within one specific group.

Question 3

Using the **training** data set, create a correlation matrix of all continuous variables. Visualize the matrix and describe any patterns you see. Are any predictors correlated with each other? Which ones, and in which direction?

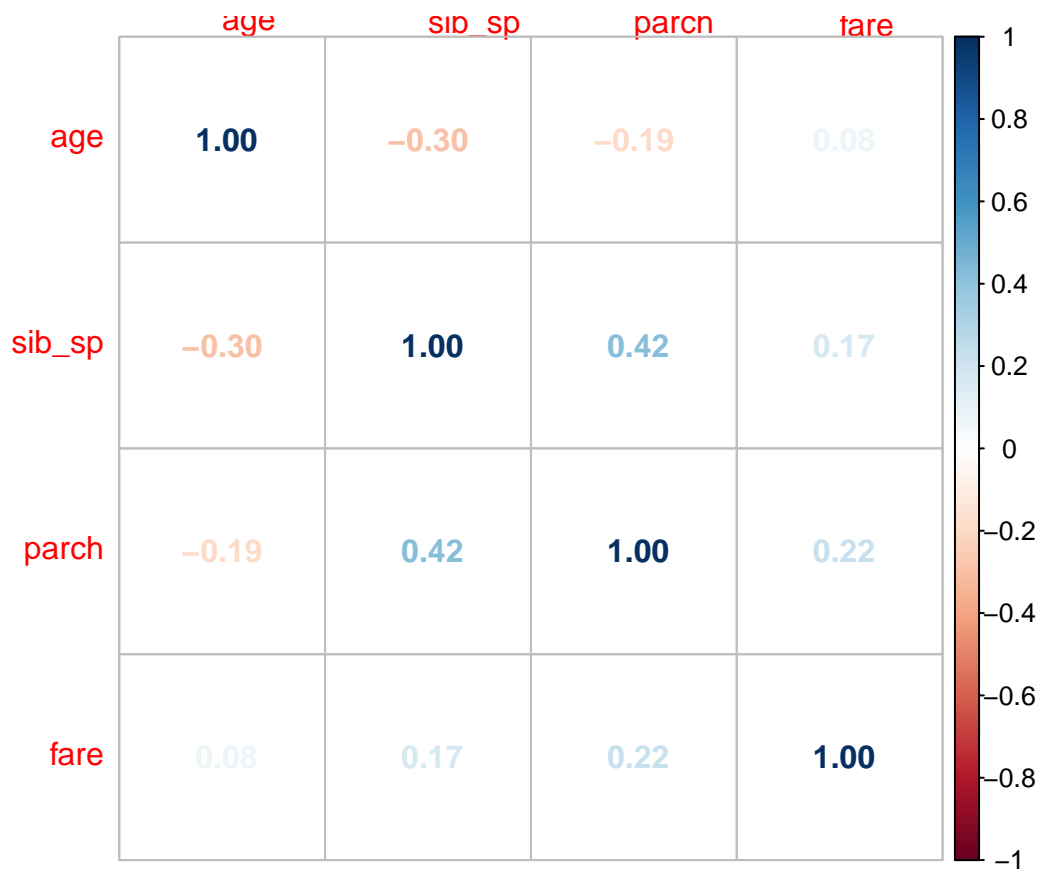
```

# Select continuous variables
continuous_vars <- select(train_data, age, sib_sp, parch, fare)

# Compute the correlation matrix
cor_matrix <- cor(contiguous_vars, use = "pairwise.complete.obs")

# Visualize the correlation matrix (with rotated x-axis labels to read easierly)
corrplot(cor_matrix, method = "number", tl.srt = 360)

```



Negative Correlation between sib_sp and age: In this context, it suggests that younger passengers are more likely to be traveling with siblings or spouses, which does make sense.

Positive Correlation between sib_sp and parch: This indicates that passengers who are traveling with siblings or spouses are also more likely to be traveling with parents or children. Again, this makes sense as families are likely to travel together.

Also, 'Fare' is not strongly correlated with other variables, probably because it is possibly more related to social or economic status, which isn't explained by the other continuous variables like 'parch' or 'sib_sp'.

Question 4

Using the **training** data, create a recipe predicting the outcome variable **survived**. Include the following predictors: ticket class, sex, age, number of siblings or spouses aboard, number of parents or children aboard, and passenger fare.

Recall that there were missing values for **age**. To deal with this, add an imputation step using `step_impute_linear()`. Next, use `step_dummy()` to **dummy** encode categorical predictors. Finally, include interactions between:

- Sex and passenger fare, and
- Age and passenger fare.

You'll need to investigate the `tidymodels` documentation to find the appropriate step functions to use.

```
# Define the recipe for Question 4
recipe_q4 <- recipe(survived ~ pclass + sex + age + sib_sp + parch + fare,
                    data = train_data) %>%
  # Impute missing values for 'age' using linear imputation
  step_impute_linear(age, impute_with= imp_vars(sib_sp)) %>%
  # Dummy encode categorical predictors
  step_dummy(all_nominal_predictors()) %>%
  # Include interaction terms
  step_interact(~ starts_with("sex") : age + age:fare)

# Print and check the recipe
recipe_q4
```

Question 5

Specify a **logistic regression** model for classification using the "glm" engine. Then create a workflow. Add your model and the appropriate recipe. Finally, use `fit()` to apply your workflow to the **training** data.

Hint: Make sure to store the results of `fit()`. You'll need them later on.

```
# Specify a logistic regression model with the 'glm' engine
logistic_spec <- logistic_reg() %>%
  set_engine("glm") %>%
  set_mode("classification")

# Create a workflow combining the recipe and model specification
workflow_q5 <- workflow() %>%
  add_recipe(recipe_q4) %>%
  add_model(logistic_spec)

# Store the results of fit()
fit_workflow_q5 <- fit(workflow_q5, data = train_data)

# Display the results to check the model fit
fit_workflow_q5
```

```
## == Workflow [trained] =====
## Preprocessor: Recipe
```

```
## Model: logistic_reg()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_impute_linear()
## * step_dummy()
## * step_interact()
##
## -- Model -----
##
## Call: stats::glm(formula = .y ~ ., family = stats::binomial, data = data)
##
## Coefficients:
##      (Intercept)          age      sib_sp      parch          fare
##      3.1618660     -0.0153437     -0.3572281     -0.2346975     -0.0099770
##      pclass_X2      pclass_X3      sex_male  sex_male_x_age      age_x_fare
##      -0.7401441     -2.0048715     -1.0795435     -0.0641477      0.0004169
##
## Degrees of Freedom: 622 Total (i.e. Null);  613 Residual
## Null Deviance:      829.6
## Residual Deviance: 539.7      AIC: 559.7
```

Question 6

Repeat Question 5, but this time specify a linear discriminant analysis model for classification using the "MASS" engine.

```
# Specify a Linear Discriminant Analysis model with the 'MASS' engine
lda_spec <- discrim_linear() %>%
  set_engine("MASS") %>%
  set_mode("classification")

# Create a new workflow combining recipe_q4 and the new LDA model specification.
workflow_q6 <- workflow() %>%
  add_recipe(recipe_q4) %>%
  add_model(lda_spec)

# Store the results of fit() in a new variable
fit_workflow_q6 <- fit(workflow_q6, data = train_data)

# Display the results to check the model fit
fit_workflow_q6
```

```
## == Workflow [trained] =====
## Preprocessor: Recipe
## Model: discrim_linear()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_impute_linear()
## * step_dummy()
```

```
## * step_interact()
##
## -- Model -----
## Call:
## lda(..y ~ ., data = data)
##
## Prior probabilities of groups:
##      No      Yes
## 0.6163724 0.3836276
##
## Group means:
##      age      sib_sp      parch      fare pclass_X2 pclass_X3 sex_male
## No  29.68856 0.5677083 0.3567708 22.48433 0.1744792 0.6848958 0.8411458
## Yes 28.31317 0.4435146 0.4853556 46.90032 0.2887029 0.3640167 0.2970711
##      sex_male_x_age age_x_fare
## No      25.908660    665.0898
## Yes      7.923317   1439.8651
##
## Coefficients of linear discriminants:
##                      LD1
## age                -0.0045217512
## sib_sp              -0.2312120167
## parch               -0.1410511932
## fare                -0.0024362813
## pclass_X2           -0.3585050280
## pclass_X3           -1.2124231243
## sex_male            -1.1837871204
## sex_male_x_age      -0.0357868566
## age_x_fare          0.0001345862
```

Question 7

Repeat Question 5, but this time specify a quadratic discriminant analysis model for classification using the "MASS" engine.

```
# Create a parsnip model specification for QDA
qda_spec <- discrim_quad() %>%
  set_engine('MASS') %>%
  set_mode('classification')

# Create a new workflow combining recipe_q4 and the new QDA model specification.
workflow_q7 <- workflow() %>%
  add_recipe(recipe_q4) %>%
  add_model(qda_spec)

# Store the results of fit() in a new variable
fit_workflow_q7 <- fit(workflow_q7, data = train_data)

# Display the results to check the model fit
fit_workflow_q7
```

```
## == Workflow [trained] =====
## Preprocessor: Recipe
```



```
## Model: discrim_quad()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_impute_linear()
## * step_dummy()
## * step_interact()
##
## -- Model -----
## Call:
## qda(..y ~ ., data = data)
##
## Prior probabilities of groups:
##      No      Yes
## 0.6163724 0.3836276
##
## Group means:
##      age      sib_sp      parch      fare pclass_X2 pclass_X3 sex_male
## No  29.68856 0.5677083 0.3567708 22.48433 0.1744792 0.6848958 0.8411458
## Yes 28.31317 0.4435146 0.4853556 46.90032 0.2887029 0.3640167 0.2970711
##      sex_male_x_age age_x_fare
## No      25.908660    665.0898
## Yes      7.923317    1439.8651
```

Question 8

Repeat Question 5, but this time specify a k -nearest neighbors model for classification using the "kkn" engine. Choose a value for k to try.

```
# Specify a k-NN model with the 'kkn' engine and choose k (e.g., k = 8)
knn_spec <- nearest_neighbor(neighbors = 8) %>%
  set_engine("kkn") %>%
  set_mode("classification")

# Create a new workflow combining recipe_q4 and the new k nearest neighbors model
workflow_q8 <- workflow() %>%
  add_recipe(recipe_q4) %>%
  add_model(knn_spec)

# Store the results of fit() in a new variable
fit_workflow_q8 <- fit(workflow_q8, data = train_data)

# Display the results to check the model fit
fit_workflow_q8
```

```
## == Workflow [trained] =====
## Preprocessor: Recipe
## Model: nearest_neighbor()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
```

```
## * step_impute_linear()
## * step_dummy()
## * step_interact()
##
## -- Model -----
##
## Call:
## kkn::train.kknn(formula = ..y ~ ., data = data, ks = min_rows(8,      data, 5))
##
## Type of response variable: nominal
## Minimal misclassification: 0.1701445
## Best kernel: optimal
## Best k: 8
```

Question 9

Now you've fit four different models to your training data.

Use `predict()` and `bind_cols()` to generate predictions using each of these 4 models and your **training** data. Then use the metric of **area under the ROC curve** to assess the performance of each of the four models.

```
# Generate predictions for each of the 4 models using the training data
pred_logistic <- predict(fit_workflow_q5, new_data = train_data,
                        type = "prob") %>% as_tibble()
pred_lda <- predict(fit_workflow_q6, new_data = train_data,
                  type = "prob") %>% as_tibble()
pred_qda <- predict(fit_workflow_q7, new_data = train_data,
                  type = "prob") %>% as_tibble()
pred_knn <- predict(fit_workflow_q8, new_data = train_data,
                  type = "prob") %>% as_tibble()

# Create data frames to store truth and predicted probabilities
df_logistic <- tibble(truth = as.factor(train_data$survived),
                    .pred_Yes = pred_logistic$.pred_Yes)
df_lda <- tibble(truth = as.factor(train_data$survived),
                .pred_Yes = pred_lda$.pred_Yes)
df_qda <- tibble(truth = as.factor(train_data$survived),
                .pred_Yes = pred_qda$.pred_Yes)
df_knn <- tibble(truth = as.factor(train_data$survived),
                .pred_Yes = pred_knn$.pred_Yes)

# Reorder the factor levels so that "Yes" is the reference level for each
df_logistic$truth <- relevel(df_logistic$truth, ref = "Yes")
df_lda$truth <- relevel(df_lda$truth, ref = "Yes")
df_qda$truth <- relevel(df_qda$truth, ref = "Yes")
df_knn$truth <- relevel(df_knn$truth, ref = "Yes")

# Calculate AUC-ROC for each model
roc_logistic <- roc_auc(df_logistic, truth, .pred_Yes)
roc_lda <- roc_auc(df_lda, truth, .pred_Yes)
roc_qda <- roc_auc(df_qda, truth, .pred_Yes)
roc_knn <- roc_auc(df_knn, truth, .pred_Yes)
```

```
# Display AUC-ROC for each model
print(paste("AUC-ROC for Logistic Regression: ", roc_logistic$.estimate))
```

```
## [1] "AUC-ROC for Logistic Regression: 0.861118375174337"
```

```
print(paste("AUC-ROC for LDA: ", roc_lda$.estimate))
```

```
## [1] "AUC-ROC for LDA: 0.855604951185495"
```

```
print(paste("AUC-ROC for QDA: ", roc_qda$.estimate))
```

```
## [1] "AUC-ROC for QDA: 0.846626569037658"
```

```
print(paste("AUC-ROC for k-NN: ", roc_knn$.estimate))
```

```
## [1] "AUC-ROC for k-NN: 0.974301560320781"
```

Question 10

Fit all four models to your **testing** data and report the AUC of each model on the **testing** data. Which model achieved the highest AUC on the **testing** data?

```
# Generate predictions for each of the 4 models using the testing data
pred_logistic_test <- predict(fit_workflow_q5, new_data = test_data,
                             type = "prob") %>% as_tibble()
pred_lda_test <- predict(fit_workflow_q6, new_data = test_data,
                        type = "prob") %>% as_tibble()
pred_qda_test <- predict(fit_workflow_q7, new_data = test_data,
                        type = "prob") %>% as_tibble()
pred_knn_test <- predict(fit_workflow_q8, new_data = test_data,
                        type = "prob") %>% as_tibble()

# Create data frames to store truth and predicted probabilities for testing data
df_logistic_test <- tibble(truth = as.factor(test_data$survived),
                          .pred_Yes = pred_logistic_test$.pred_Yes)
df_lda_test <- tibble(truth = as.factor(test_data$survived),
                     .pred_Yes = pred_lda_test$.pred_Yes)
df_qda_test <- tibble(truth = as.factor(test_data$survived),
                     .pred_Yes = pred_qda_test$.pred_Yes)
df_knn_test <- tibble(truth = as.factor(test_data$survived),
                     .pred_Yes = pred_knn_test$.pred_Yes)

# Reorder the factor levels so that "Yes" is the reference level for each
df_logistic_test$truth <- relevel(df_logistic_test$truth, ref = "Yes")
df_lda_test$truth <- relevel(df_lda_test$truth, ref = "Yes")
df_qda_test$truth <- relevel(df_qda_test$truth, ref = "Yes")
df_knn_test$truth <- relevel(df_knn_test$truth, ref = "Yes")

# Calculate AUC-ROC for each model on testing data
roc_logistic_test <- roc_auc(df_logistic_test, truth, .pred_Yes)
```

```
roc_lda_test <- roc_auc(df_lda_test, truth, .pred_Yes)
roc_qda_test <- roc_auc(df_qda_test, truth, .pred_Yes)
roc_knn_test <- roc_auc(df_knn_test, truth, .pred_Yes)
```

```
# Display AUC-ROC for each model on testing data
print(paste("Test AUC-ROC for Logistic Regression: ",
            roc_logistic_test$.estimate))
```

```
## [1] "Test AUC-ROC for Logistic Regression: 0.864666078258312"
```

```
print(paste("Test AUC-ROC for LDA: ",
            roc_lda_test$.estimate))
```

```
## [1] "Test AUC-ROC for LDA: 0.861900558987938"
```

```
print(paste("Test AUC-ROC for QDA: ",
            roc_qda_test$.estimate))
```

```
## [1] "Test AUC-ROC for QDA: 0.844895557516918"
```

```
print(paste("Test AUC-ROC for k-NN: ",
            roc_knn_test$.estimate))
```

```
## [1] "Test AUC-ROC for k-NN: 0.851220947337452"
```

Using your top-performing model, create a confusion matrix and visualize it. Create a plot of its ROC curve.

Based on the output above, Logistic Regression model seems has the highest AUC-ROC performance.


```
# Create a new column in pred_logistic_test for predicted class labels
pred_logistic_test$.pred_class <- ifelse(pred_logistic_test$.pred_Yes > 0.5,
                                         "Yes", "No")
```

```
# Create a new tibble combining 'survived' and '.pred_class'
combined_data <- tibble(
  truth = test_data$survived,
  .pred_class = pred_logistic_test$.pred_class
)
```

```
# Convert .pred_class to a factor
combined_data$.pred_class <- as.factor(combined_data$.pred_class)
```

```
# Create the confusion matrix
confusion <- conf_mat(combined_data, truth = truth, estimate = .pred_class)
```

```
# Visualize the confusion matrix
autoplot(confusion)
```




```
# Create a new tibble combining 'survived' and predicted probabilities
roc_data <- tibble(
  truth = as.factor(test_data$survived),
  .pred_Yes = pred_logistic_test$.pred_Yes
)
```

```
# Create the ROC curve
roc_curve_data <- roc_curve(roc_data, truth, .pred_Yes)
```

```
# Plot the ROC curve using ggplot2
ggplot(roc_curve_data, aes(x = specificity, y = sensitivity)) +
  geom_line() +
  geom_abline(linetype = "dashed") +
  labs(title = "ROC Curve for Logistic Regression",
       x = "Specificity",
       y = "Sensitivity")
```

ROC Curve for Logistic Regression



How did your best model perform? Compare its **training** and **testing** AUC values. If the values differ, why do you think this is so?

```
## The best-performing model in the training data is k-NN, and the AUC_ROC value
#is as high as 0.974. While in the testing data, the best model is the Logistic
#Regression with value of approximately 0.865.
```

```
## The AUC-ROC values for both the training and testing datasets are very close
#for Logistic Regression, LDA, and QDA, suggesting good generalization. However,
#for k-NN, the AUC-ROC is significantly higher on the training data compared to
#the testing data (0.974 vs. 0.851), which could be a sign of overfitting to the
#training data.
```

```
## The Logistic Regression model, with AUC-ROC values of 0.861 on the training
#data and 0.865 on the test data, appears to be the most reliable and consistent
#model for this specific project. The small difference between the training and
#testing AUC values suggests that the model has generalized well to the unseen
#data, without signs of overfitting or underfitting.
```