

Homework 4

PSTAT Winter 2023

Haocheng Zhang

2023-03-17

```
library(ISLR)
data(Auto)
```

```
## Q1(a):
# Qualitative predictors: cylinders, year, origin, (name, but it is excluded
# from the linear regression model)

# Quantitative predictors: displacement, horsepower, weight, acceleration
```

```
## Q1(b) Fit the MLR model.
lmod<- lm(mpg~. - name, Auto)
# Then use the anova() function in R to perform an analysis of variance
anova_lmod <- anova(lmod)

print(anova_lmod)
```

```
## Analysis of Variance Table
##
## Response: mpg
##      Df Sum Sq Mean Sq F value    Pr(>F)
## cylinders    4 15274.5   3818.6  470.9373 < 2.2e-16 ***
## displacement  1  1098.0   1098.0  135.4069 < 2.2e-16 ***
## horsepower    1   588.0    588.0   72.5161  4.25e-16 ***
## weight        1   715.1    715.1   88.1961 < 2.2e-16 ***
## acceleration  1     7.7     7.7    0.9457   0.3315
## year         12  2960.4    246.7   30.4251 < 2.2e-16 ***
## origin        2   183.2     91.6   11.2972  1.73e-05 ***
## Residuals    369  2992.1     8.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Based on the anova table and p-values, we can make the following observations about the
# predictors and their linear association with 'mpg' conditional on the other
# predictors in the model:
```

```
# cylinders: p-value < 2.2e-16<0.05, reject the null hypothesis
# displacement: p-value < 2.2e-16<0.05, reject the null hypothesis
# horsepower: p-value = 4.25e-16<0.05, reject the null hypothesis
# weight: p-value < 2.2e-16<0.05, reject the null hypothesis
```

```
# acceleration: p-value = 0.3315>0.05, fail to reject the null hypothesis
# year: p-value < 2.2e-16<0.05, reject the null hypothesis
# origin: p-value = 1.73e-05<0.05, reject the null hypothesis
```

```
# Q1(c) To predict the mpg with given specifications, we use predict() function.
```

```
# Create a new data frame with the desired values for each predictor:
```

```
Japanese_car <- data.frame(
  name = factor("Jpcar", levels = levels(Autod$name)),
  cylinders = factor(3, levels = levels(Autod$cylinders)),
  displacement = 100,
  horsepower = 85,
  weight = 3000,
  acceleration = 20,
  year = factor(80, levels = levels(Autod$year)),
  origin = factor(3, levels = levels(Autod$origin))
)
```

```
# Then, use the predict() function with the fitted MLR model
```

```
predicted_mpg <- predict(lmod, newdata = Japanese_car)
```

```
# Print the predicted mpg
```

```
show(predicted_mpg)
```

```
##          1
```

```
## 24.64804
```

```
# Q1(d) We may extract the coefficients from the fitted MLR model,
```

```
# and compute the difference
```

```
diff_origin <- coef(lmod)["origin3"] - coef(lmod)["origin2"]
```

```
# Print the difference between Japanese cars and European cars.
```

```
show(diff_origin)
```

```
##   origin3
```

```
## 0.5996414
```

```
# Q1(e) To fit a model to predict mpg using origin, horsepower,
```

```
# and their interaction
```

```
interaction_model <- lm(mpg ~ origin * horsepower, data = Autod)
```

```
# Print the summary of the interaction_model
```

```
summary(interaction_model)
```

```
##
```

```
## Call:
```

```
## lm(formula = mpg ~ origin * horsepower, data = Autod)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -10.7415  -2.9547  -0.6389   2.3978  14.2495
```

```
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    34.476496   0.890665  38.709 < 2e-16 ***
## origin2        10.997230   2.396209   4.589 6.02e-06 ***
## origin3        14.339718   2.464293   5.819 1.24e-08 ***
## horsepower     -0.121320   0.007095 -17.099 < 2e-16 ***
## origin2:horsepower -0.100515  0.027723  -3.626 0.000327 ***
## origin3:horsepower -0.108723  0.028980  -3.752 0.000203 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.422 on 386 degrees of freedom
## Multiple R-squared:  0.6831, Adjusted R-squared:  0.679
## F-statistic: 166.4 on 5 and 386 DF, p-value: < 2.2e-16
```

```
# Based on the summary output of the fitted model, the fitted linear model
# using origin, horsepower, and their interaction can be written as:
# mpg = 34.4765 + 10.9972 * origin2 + 14.3397 * origin3 - 0.1213 * horsepower
#       - 0.1005 * (origin2 * horsepower) - 0.1087 * (origin3 * horsepower)
#       + $\\epsilon$
# Here, the variables are:
# mpg: miles per gallon
# origin2: 1 if the car is European, 0 otherwise
# origin3: 1 if the car is Japanese, 0 otherwise
# horsepower: engine horsepower
# $\\epsilon$: random error term
```

```
# Q1(f) use the Akaike Information Criterion (AIC) to compare different degrees
# of polynomial regression models and choose the one with the lowest AIC value.
```

```
# Maximum polynomial degree to test
max_degree <- 5

# Create an empty vector to store AIC values
aic_values <- numeric(max_degree)

# Loop through different polynomial degrees and fit models
for (degree in 1:max_degree) {
  poly_model <- lm(mpg ~ poly(weight, degree, raw = TRUE), data = Autod)
  aic_values[degree] <- AIC(poly_model)
}

# Find the degree with the lowest AIC value
best_degree <- which.min(aic_values)

# Print the AIC values and the best degree
print(aic_values)
```

```
## [1] 2265.939 2238.115 2240.113 2241.657 2242.807
```

```
print(best_degree)
```

```
## [1] 2
```

```
# From the outcome above, we know 2nd degree is a proper degree of polynomial,  
# because it has the lowest AIC value.
```

```
# Q1(g) Load the MASS package for further use.  
library(MASS)
```

```
# Full model with all predictors (except name)  
full_model <- lm(mpg ~ . - name, data = Autod)
```

```
# Perform a backward selection using 'stepAIC()' function  
best_model <- stepAIC(full_model, direction = "backward")
```

```
## Start: AIC=842.72  
## mpg ~ (cylinders + displacement + horsepower + weight + acceleration +  
## year + origin + name) - name  
##
```

| | Df | Sum of Sq | RSS | AIC |
|-------------------|----|-----------|--------|---------|
| ## - acceleration | 1 | 0.01 | 2992.1 | 840.72 |
| ## <none> | | | 2992.1 | 842.72 |
| ## - displacement | 1 | 24.70 | 3016.8 | 843.94 |
| ## - horsepower | 1 | 73.45 | 3065.5 | 850.23 |
| ## - origin | 2 | 183.21 | 3175.3 | 862.02 |
| ## - cylinders | 4 | 472.77 | 3464.8 | 892.23 |
| ## - weight | 1 | 558.60 | 3550.7 | 907.82 |
| ## - year | 12 | 2831.60 | 5823.7 | 1079.78 |

```
##  
## Step: AIC=840.72  
## mpg ~ cylinders + displacement + horsepower + weight + year +  
## origin  
##
```

| | Df | Sum of Sq | RSS | AIC |
|-------------------|----|-----------|--------|---------|
| ## <none> | | | 2992.1 | 840.72 |
| ## - displacement | 1 | 24.88 | 3017.0 | 841.97 |
| ## - horsepower | 1 | 115.58 | 3107.7 | 853.58 |
| ## - origin | 2 | 183.45 | 3175.5 | 860.05 |
| ## - cylinders | 4 | 476.39 | 3468.5 | 890.64 |
| ## - weight | 1 | 730.02 | 3722.1 | 924.31 |
| ## - year | 12 | 2841.52 | 5833.6 | 1078.45 |

```
# Print the summary of the best model  
summary(best_model)
```

```
##  
## Call:  
## lm(formula = mpg ~ cylinders + displacement + horsepower + weight +  
## year + origin, data = Autod)  
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.931 -1.671 -0.049  1.448 11.612
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  30.9706782  1.9703396  15.718 < 2e-16 ***
## cylinders4    6.9489835  1.5189655   4.575 6.51e-06 ***
## cylinders5    6.6467365  2.3240427   2.860 0.004477 **
## cylinders6    4.3050676  1.6932440   2.542 0.011413 *
## cylinders8    6.3723261  1.9615936   3.249 0.001266 **
## displacement  0.0117929  0.0067234   1.754 0.080256 .
## horsepower   -0.0395543  0.0104627  -3.781 0.000182 ***
## weight       -0.0051675  0.0005439  -9.501 < 2e-16 ***
## year71        0.9057500  0.8066633   1.123 0.262236
## year72       -0.4921367  0.8015260  -0.614 0.539593
## year73       -0.5550656  0.7185757  -0.772 0.440340
## year74        1.2376123  0.8470486   1.461 0.144840
## year75        0.8654150  0.8276285   1.046 0.296402
## year76        1.4923994  0.7942429   1.879 0.061027 .
## year77        2.9948793  0.8136242   3.681 0.000267 ***
## year78        2.9703034  0.7736654   3.839 0.000145 ***
## year79        4.8922614  0.8182998   5.979 5.30e-09 ***
## year80        9.0552685  0.8695618  10.414 < 2e-16 ***
## year81        6.4527050  0.8525066   7.569 3.02e-13 ***
## year82        7.8336547  0.8429257   9.293 < 2e-16 ***
## origin2       1.6931856  0.5155093   3.284 0.001119 **
## origin3       2.2936695  0.4957722   4.626 5.15e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.844 on 370 degrees of freedom
## Multiple R-squared:  0.8744, Adjusted R-squared:  0.8673
## F-statistic: 122.6 on 21 and 370 DF, p-value: < 2.2e-16
```

```
# From the output of the summary above, the best model selected by the stepAIC()
# function includes the following predictor variables: cylinders, displacement,
# horsepower, weight, year, and origin. The model has an adjusted R-squared
# value of 0.8673, which indicates a good fit. The AIC value for this model
# is 840.72, lower than the full model with all predictors, suggesting that
# it is a better model according to the AIC criterion.
```

```
# Q2(a) load the fat dataset:
```

```
library(faraway)
data(fat)
```

```
# remove every tenth observation for use as a test sample.
test_sample <- fat[seq(10, nrow(fat), by=10),]
```

```
# The remaining data will be used as a training sample for further use:
training_sample <- fat[-seq(10, nrow(fat), by=10),]
```

```
# fit a linear regression model with all predictors,
# excluding "brozek" and "density", using the training data:
```

```
training_model <- lm(siri ~ . - brozek - density, data=training_sample)

summary(training_model)
```

```
##
## Call:
## lm(formula = siri ~ . - brozek - density, data = training_sample)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.8314 -0.6722  0.1828  0.9150  6.6619
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -12.591885   6.448868  -1.953  0.052193 .
## age          0.007978   0.012320   0.648  0.517983
## weight       0.362999   0.023314  15.570 < 2e-16 ***
## height       0.049026   0.040315   1.216  0.225315
## adipos      -0.514032   0.114074  -4.506  1.09e-05 ***
## free        -0.564773   0.014889 -37.933 < 2e-16 ***
## neck         0.016525   0.089863   0.184  0.854272
## chest        0.120219   0.039590   3.037  0.002694 **
## abdom        0.140108   0.042186   3.321  0.001056 **
## hip          0.006197   0.056101   0.110  0.912148
## thigh        0.195057   0.054460   3.582  0.000424 ***
## knee         0.106637   0.093534   1.140  0.255542
## ankle        0.125118   0.081303   1.539  0.125325
## biceps       0.096199   0.064656   1.488  0.138278
## forearm      0.230775   0.073332   3.147  0.001888 **
## wrist        0.139279   0.206804   0.673  0.501378
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.55 on 211 degrees of freedom
## Multiple R-squared:  0.9692, Adjusted R-squared:  0.967
## F-statistic: 442.5 on 15 and 211 DF, p-value: < 2.2e-16
```

```
# Q2(b)
library(MASS)

# Scale all predictors, excluding 'siri', 'brozek', and 'density'
training_sample_scaled <- scale(training_sample[, !(colnames(training_sample)
%in% c("siri", "brozek", "density"))],
                                center = TRUE, scale = TRUE)
training_sample_scaled <- as.data.frame(training_sample_scaled)

# Add 'siri' back
training_sample_scaled$siri <- training_sample$siri

# Use 'lm.ridge()' function to fit a ridge regression model
rgmod <- lm.ridge(siri ~ ., data = training_sample_scaled,
                  lambda = seq(0, 100, length.out = 100))
```

```
matplot(rgmod$lambda, coef(rgmod), type="l", xlab = "lambda", ylab = "Beta hat",  
        , cex=0.8)
```

